

Lecture 3

Process Scheduling

Operating Systems

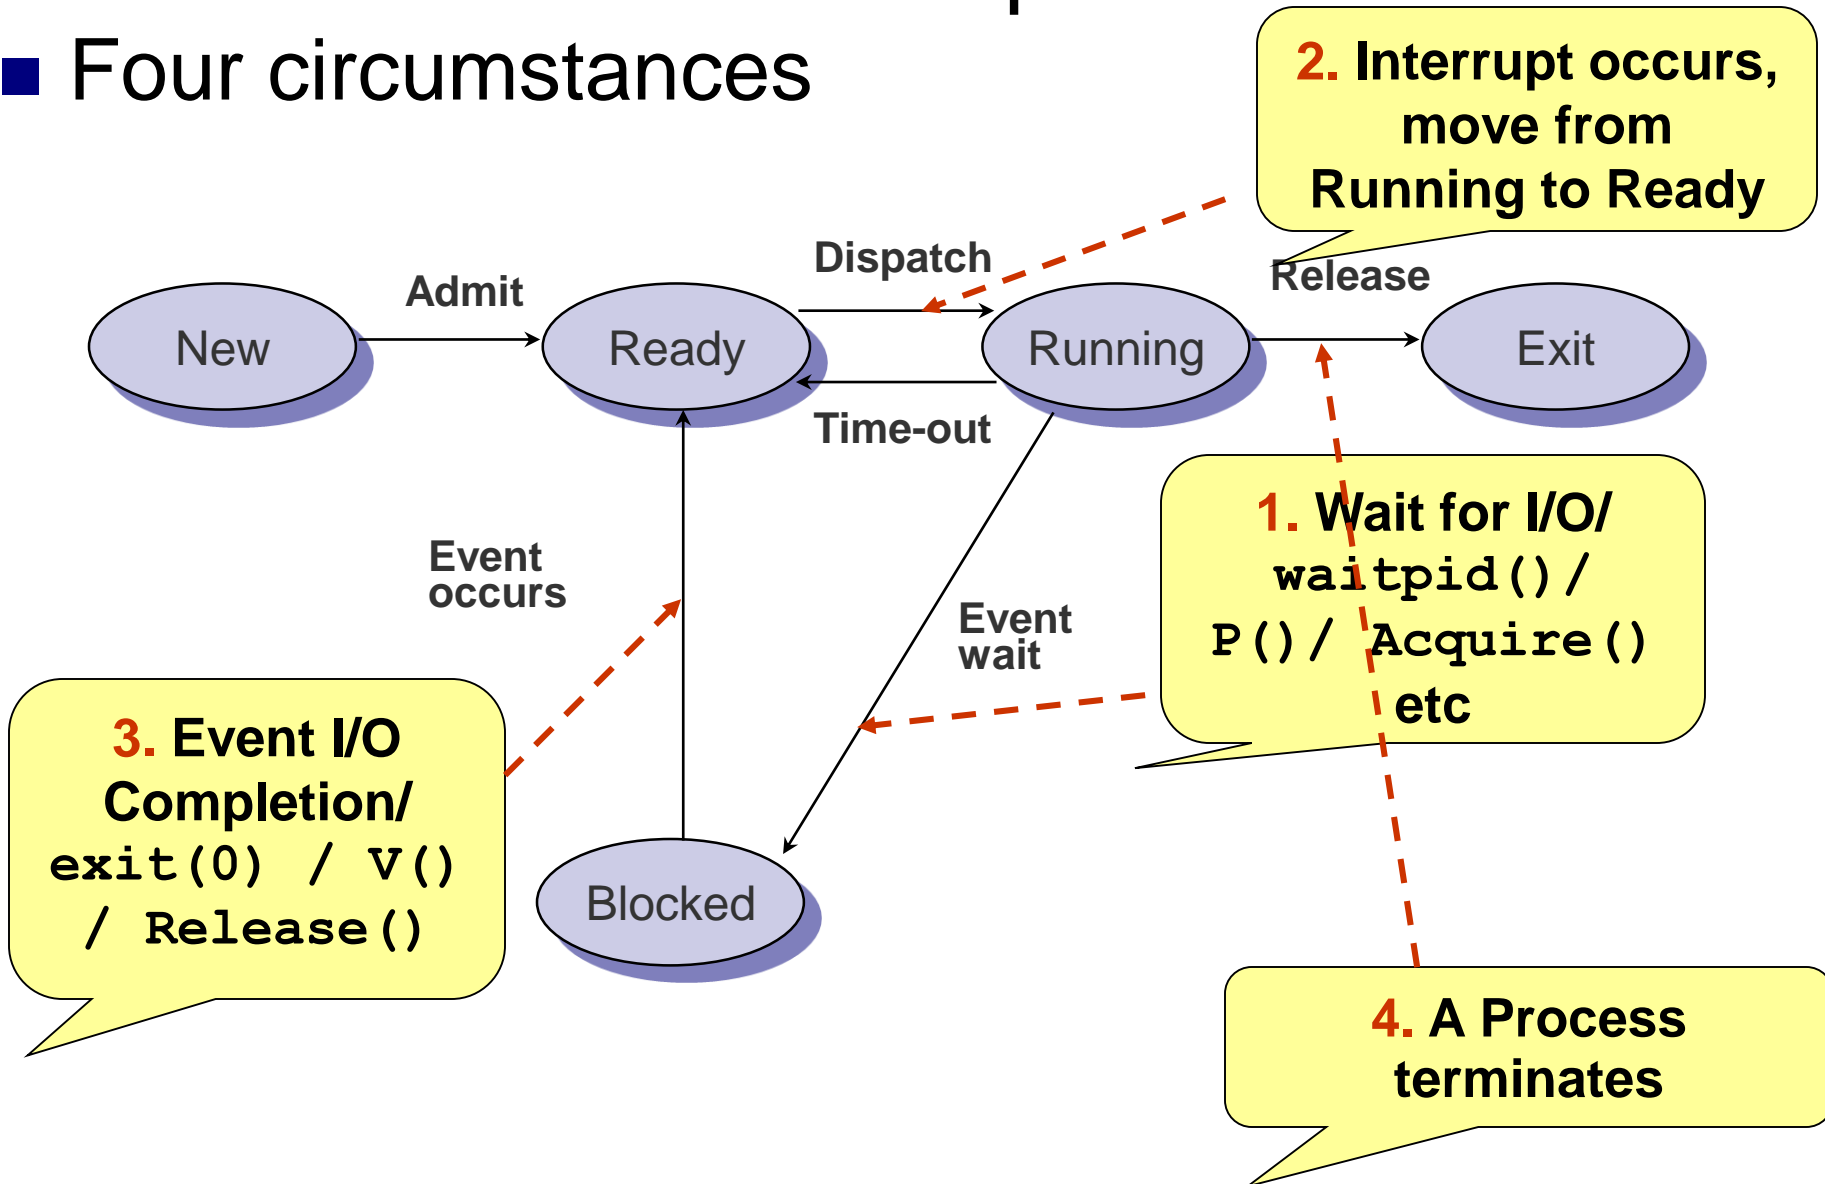


Scheduling

- Short term scheduler (CPU Scheduler)
 - Whenever the CPU becomes idle, a process must be selected for execution
 - The Process is selected from the Ready queue
- Ready queue is not necessarily a FIFO queue
- It can be
 - Priority based
 - A Tree
 - Unordered linked list etc

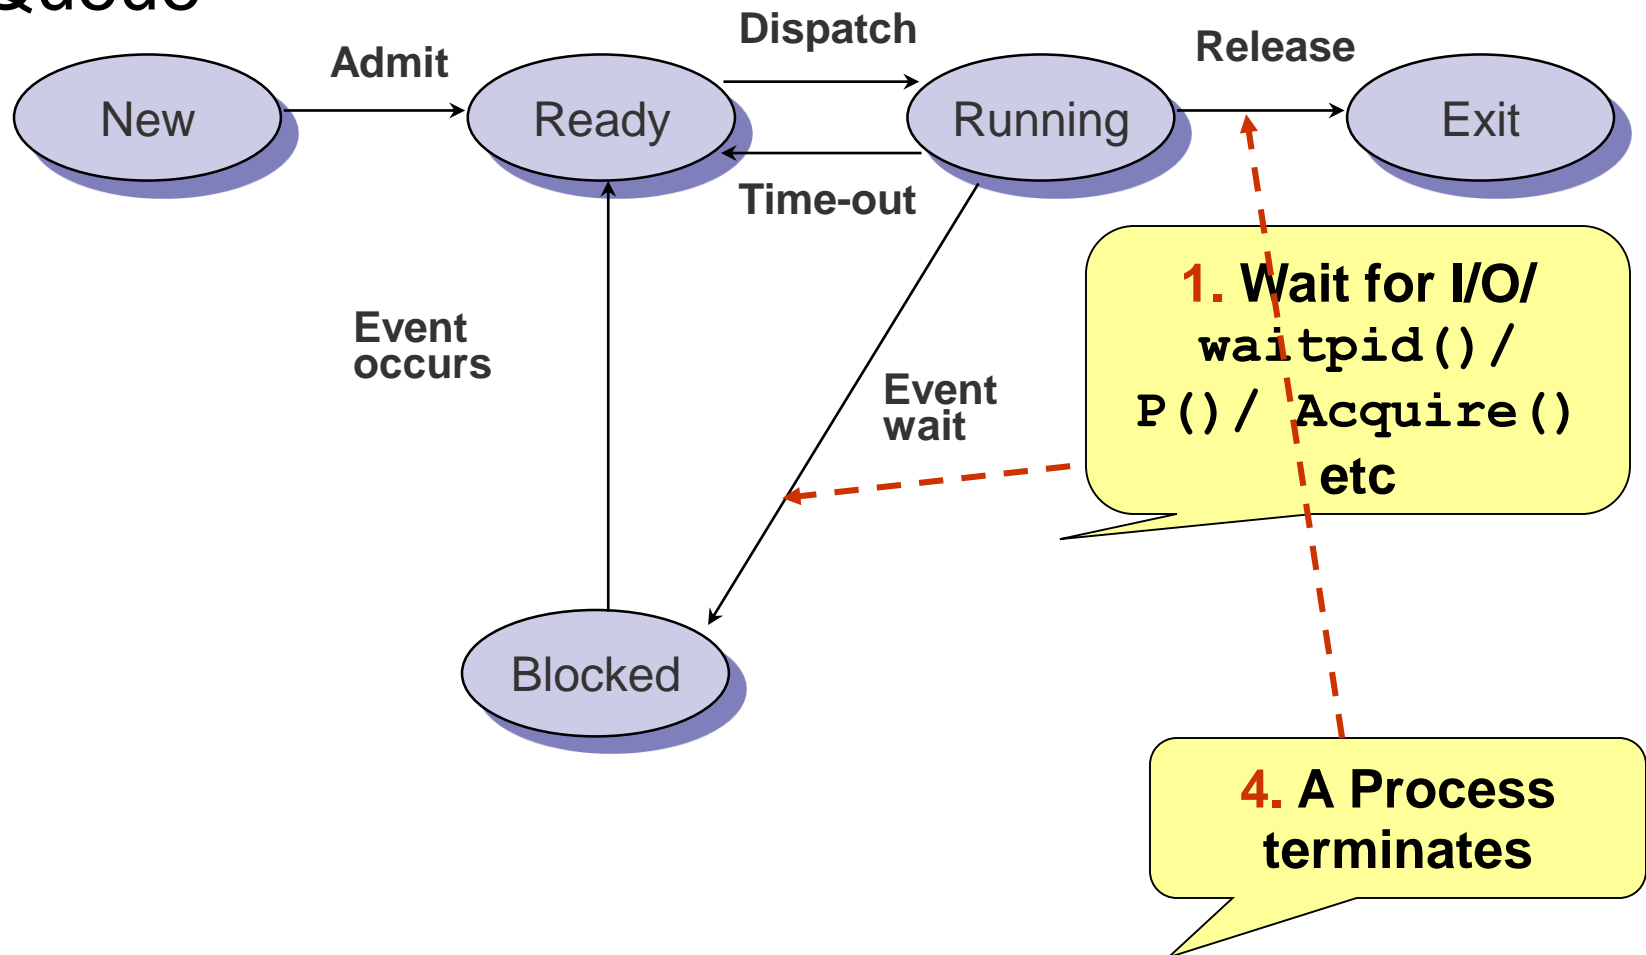
When to select a new process to Run

■ Four circumstances



Non Preemptive Scheduling

- Only the case 1 and 4
- Must select a new process, if any, from the Ready Queue

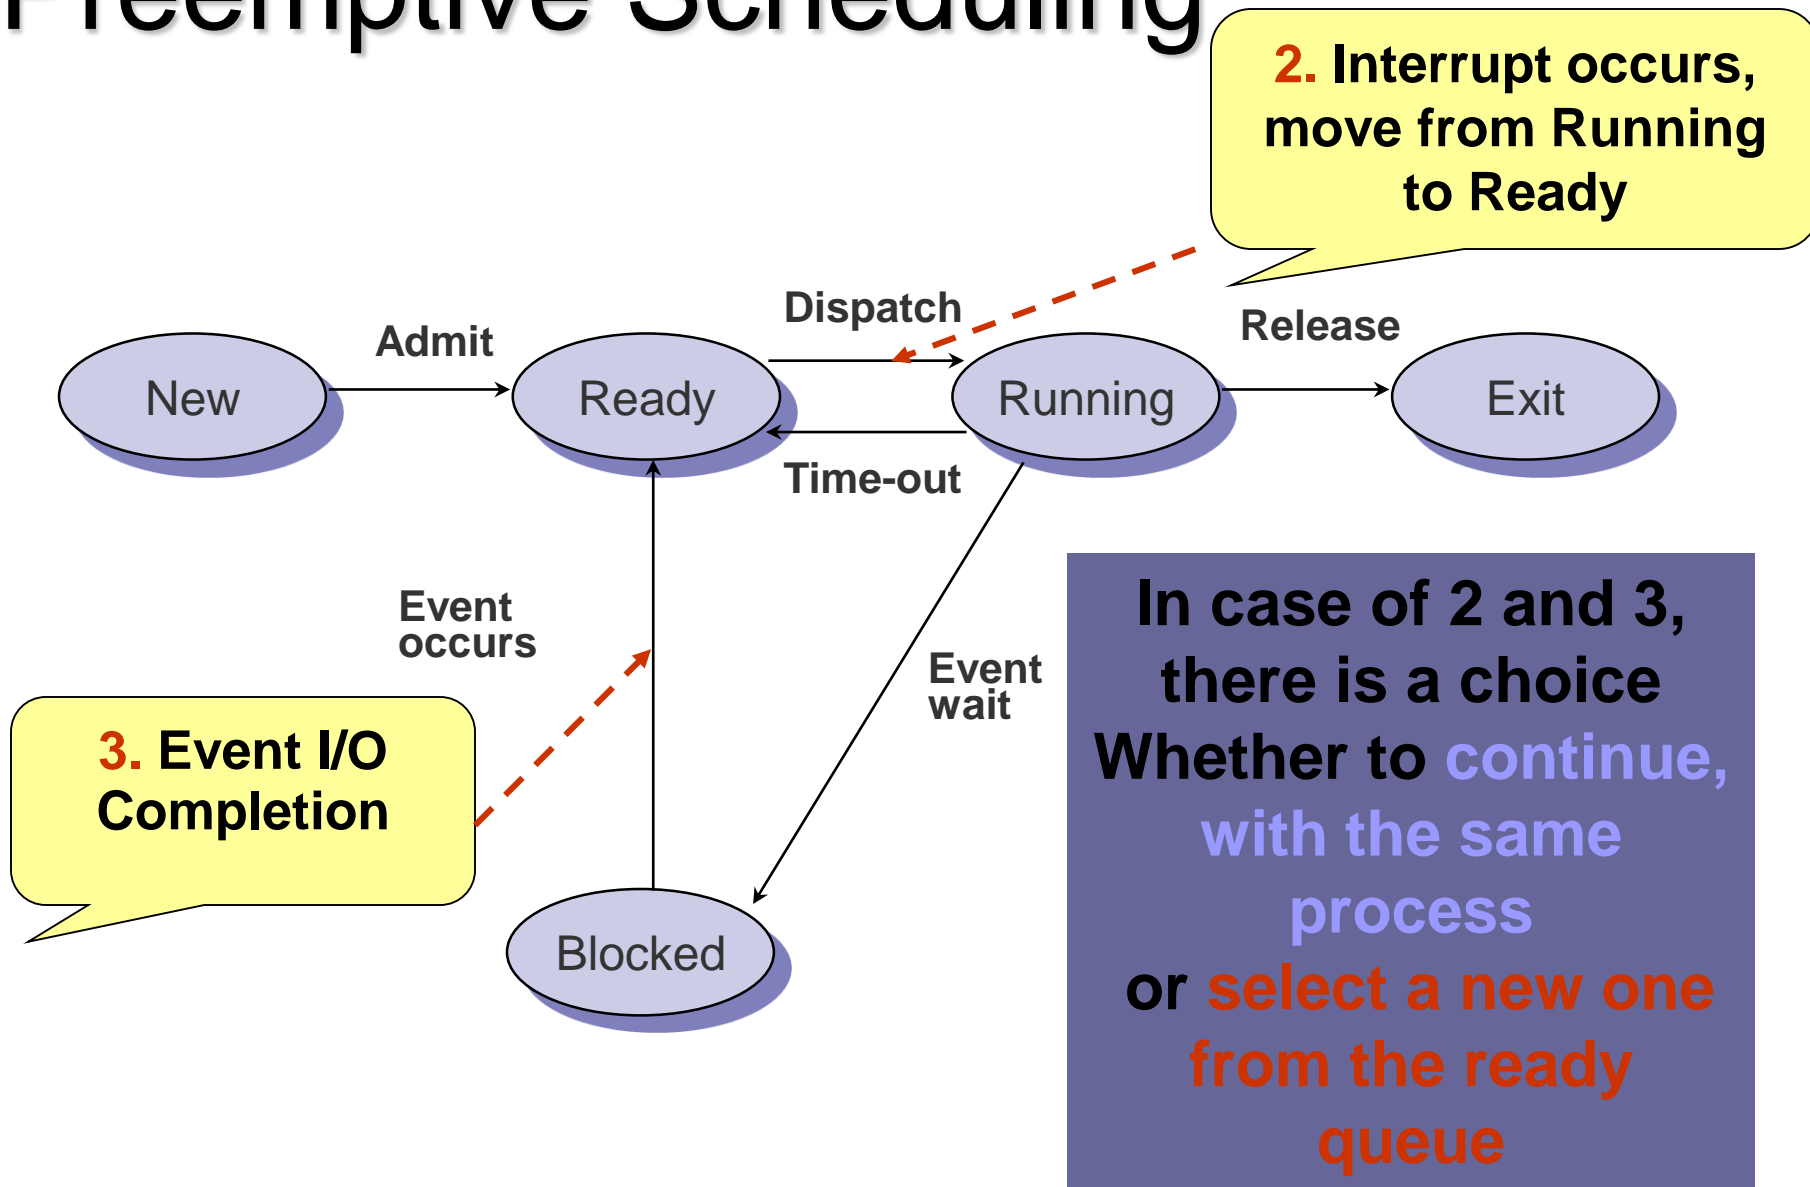




Non Preemptive Scheduling

- Once the CPU has been allocated to a process
- The process keeps it until
 - It Terminates
 - Or has to wait for:
 - I/O
 - Mutex
 - Child process
 - Semaphore
 - Conditional Variables etc
- There is no way, to get the CPU back, FORCEFULLY

Preemptive Scheduling





Scheduling Issues

- Fairness

- Don't starve process

- Priorities

- Most important first

- Deadlines

- Task X must be done by time t

- Optimization

- Throughput, response time

- Reality - No universal scheduling policy

- Many models

Optimization Criteria

■ CPU Utilization

- Keep the CPU as busy as is possible
- May range from 0% to 100%

■ Throughput

- Number of processes completed per unit time
- E.g. long processes
 - 1 process / hr
- Short processes
 - 10 processes / hr

Optimization Criteria

■ Turnaround Time

□ How long it take to execute a Process

□ $\text{Turnaround} = \text{Completion_Time}$

– Submission_Time

□ $\text{Turnaround} = \text{Wait_Time}_{\text{GetIntoMemory}}$

+ $\text{Wait_Time}_{\text{ReadyQueue}}$

+ $\text{Wait_Time}_{\text{BlockQueue}}$

+ $\text{CPU_Execution_Time}$



Optimization Criteria

- Scheduling Algorithm does not effect the waiting time in Block Queue
- It only effect the Waiting Time in the Ready Queue
- Waiting Time
 - Sum of the periods spent waiting in the Ready Queue

Optimization Criteria

- Turnaround Time is not a good criteria for Interactive Systems
- A process may
 - Produce “Some” output
 - Computes new results, while previous results are output to the user
- Response Time
- $\text{Response_Time} = \text{First_Response_Start_Time} - \text{Submission_Time}$



Optimization Criteria - Summary

- We would like to Maximize

- CPU Utilization

- Throughput

- And Minimize

- Turnaround Time

- Waiting Time

- Response Time



Scheduling Algorithms

- First come, First serve
- Shortest Job First
- Priority Scheduling
- Round-Robin Scheduling
- Multi-level Queue Scheduling
- Multi-level Feed back queue Scheduling

First come, First serve


- Simplest scheduling algorithm:
 - Run jobs in order that they arrive
- Uni-programming:
 - Run until done
- Multi-programming:
 - Run until done or Blocks on I/O
- Non-preemptive
 - A Process keeps CPU until done or I/O
- Advantage:
 - Simplicity

First come, First serve

■ Disadvantage

- Wait time depends on arrival order
- Unfair to later jobs
- (worst case: long job arrives first)

- ## ■ Three jobs (times: A=100, B=1, C=2) arrive in the order A, B, C




$$\begin{aligned}\text{Average} \\ \text{Waiting Time} &= (0 + 100 + 101) / 3 \\ &= 67\end{aligned}$$

First come, First serve

- Now if they arrive in the order B, C, A

cpu



$$\begin{aligned}\text{Average} \\ \text{Waiting Time} &= (0 + 1 + 3) / 3 \\ &= 1.33\end{aligned}$$

FCFS Convoy effect

- A CPU bound job will hold CPU until
 - Terminates
 - Or it causes an I/O burst
 - Rare occurrence, since the thread is CPU-bound
- Long periods where no I/O requests issued, and CPU held
- Result:
 - Poor I/O device utilization

FCFS Convoy effect : Example


- One CPU bound job, many I/O bound
- CPU bound runs
 - I/O jobs blocked in ready queue
 - I/O devices idle
- CPU bound blocks
 - I/O bound job(s) run, quickly block on I/O
- CPU bound runs again
- I/O of the I/O bound jobs completes
- CPU bound still runs while I/O devices idle (continues...)



Round robin (RR)

- Solution to job monopolizing CPU?
- Interrupt it.
 - Run job for some “time slice,”
 - When time is up, or it blocks
 - It moves to back of a FIFO queue
- Advantage:
 - Fair allocation of CPU across jobs
 - Low average waiting time when job lengths vary

Round robin (RR)




$$= (103 + 2 + 5) / 3$$

What is avg completion time?

- Good for Varying sized jobs
- But what about same-sized jobs?
- Assume 2 jobs of time =100 each:

Round Robin's Disadvantage



- Avg completion time?
- $(200 + 200) / 2 = 200$
- How does this compare with FCFS for same two jobs?
- $(100 + 200) / 2 = 150$



RR Time slice tradeoffs

- Performance depends on length of the timeslice
- Context switching isn't a free operation.
- If timeslice time is set too high (attempting to amortize context switch cost)
 - You get FCFS.
 - i.e. Processes will finish or block before their slice is up anyway
- If it's set too low you're spending all of your time context switching between threads.