

# Lecture 4

# Process Scheduling II

## **Operating Systems**



# Priority scheduling

- Not all jobs equal
  - So: rank them.
- Each process has a priority
  - Run highest priority ready job in system
- Priorities can be static or dynamic or both
- Among the Processes of equal priority
  - Round robin
  - FCFS

# Priority scheduling


- Priority scheduling can be Preemptive or Non-Preemptive
- When a process arrives and enters the Ready Queue
- Its priority is compared with the currently Running Process
- If Higher
  - Preemptive Scheduling
    - Run the New Thread
  - Non-Preemptive Scheduling
    - Continue running the Current Thread

# Priority scheduling

- High priority always runs over low priority.
- **Starvation**
  - A low Priority process may indefinitely wait for the CPU
- Solution: **Aging**
  - Gradually increase the Priority of processes that wait in the system for a long time.
- Which type of processes should be given Higher Priority:
  - I/O Bound???
  - CPU Bound???
- In order to keep I/O busy increase priority for jobs that often block on I/O

# Shortest Job First (SJF)


- Consider 4 jobs, a, b, c, d, run in lexical order



- The first (a) finishes at time a **time**
- The second (b) finishes at time a+b
- The third (c) finishes at time a+b+c
- The fourth (d) finishes at time a+b+c+d
- Therefore average completion =  $(a + (a + b) + (a + b + c) + (a + b + c + d)) / 4 = (4a + 3b + 2c + d) / 4$
- Minimizing this requires  $a \leq b \leq c \leq d$ .
- or Shortest Job First

# Shortest Job First (SJF)

- Run whatever job has smallest next CPU burst
- Can be pre-emptive or non-pre-emptive
- Example: same jobs (given jobs A, B, C)




$$\text{Average completion} = (1+3+103) / 3 = \sim 35$$

# Example of Non-Preemptive SJF

Process	Arrival Time	Burst Time
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

## ■ SJF (non-preemptive)




## • Average waiting time

$$= (0 + 6 + 3 + 7)/4 = 4$$

# Example of Preemptive SJF

Process	Arrival Time	Burst Time
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

## ■ SJF (preemptive)




## • Average waiting time

$$= (9 + 1 + 0 + 2)/4 = 3$$




# SJF vs. RR

- Two processes P1, P2



- RR with 100ms time slice:



- I/O idle ~90%
- SJF Offers better I/O utilization



# Shortest Job First

- The most important issue in SJF
  - Accuracy in estimation of Job length



# Multilevel Queue Scheduling

- Sometimes processes are classified into groups
- One classification can be:
  - Foreground (or Interactive) processes
  - Background (or batch) processes
- Different response time requirement
- => Different scheduling requirements
- Foreground processes usually have higher priorities

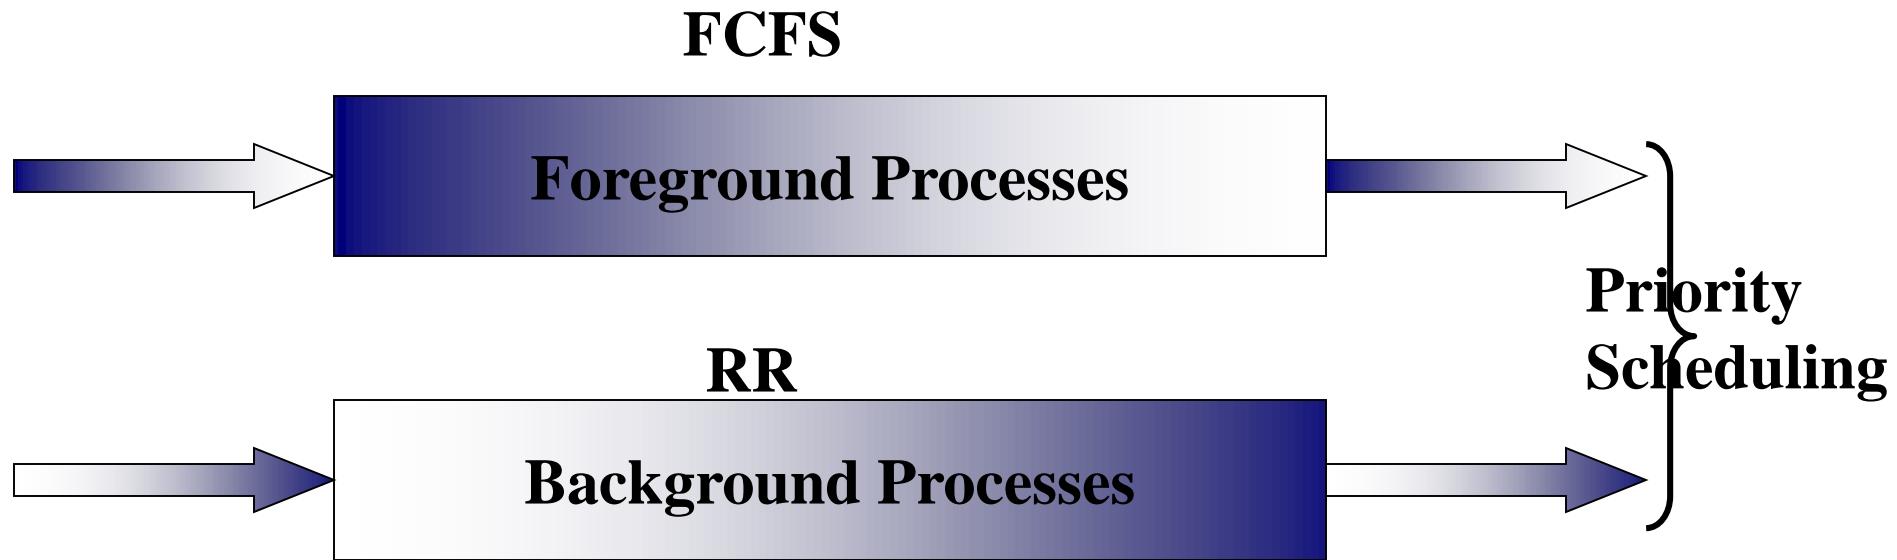


# Multilevel Queue Scheduling (MQS)

- Partition the Ready queue into a number of queues
- Processes are permanently assigned to one of the queues
- Each queue may have its own scheduling algorithm
- In addition, there must be scheduling between the queues

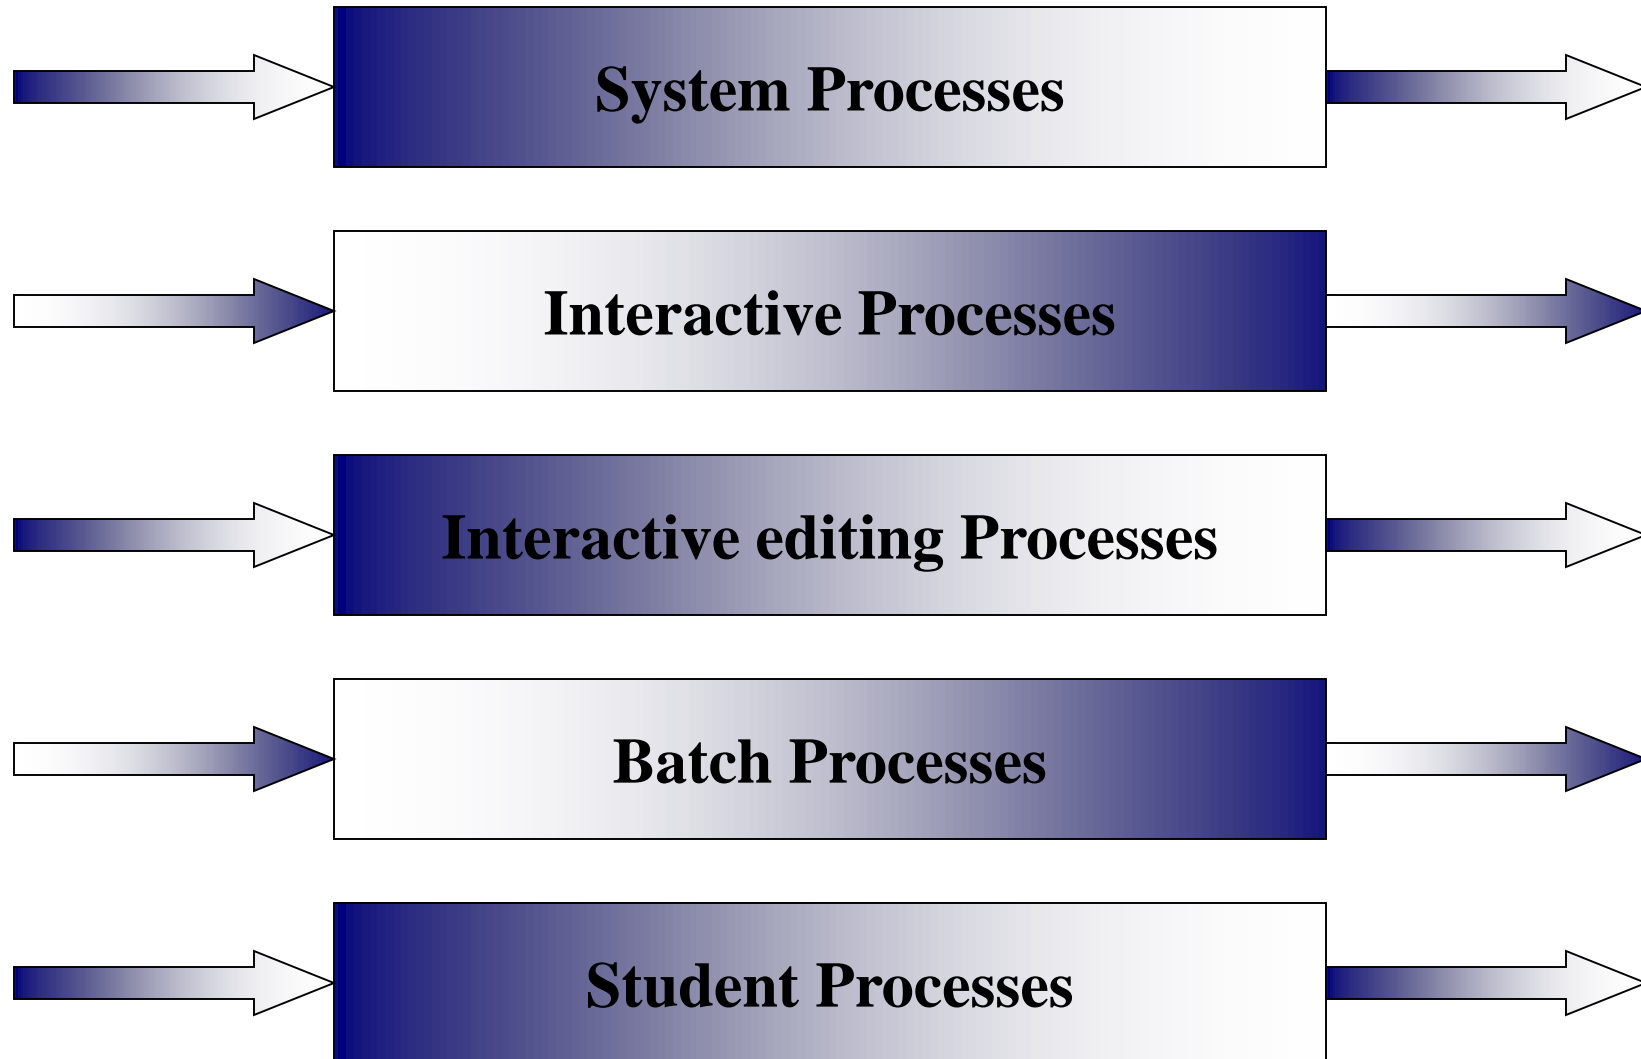
# Multilevel Queue Scheduling

- Example:



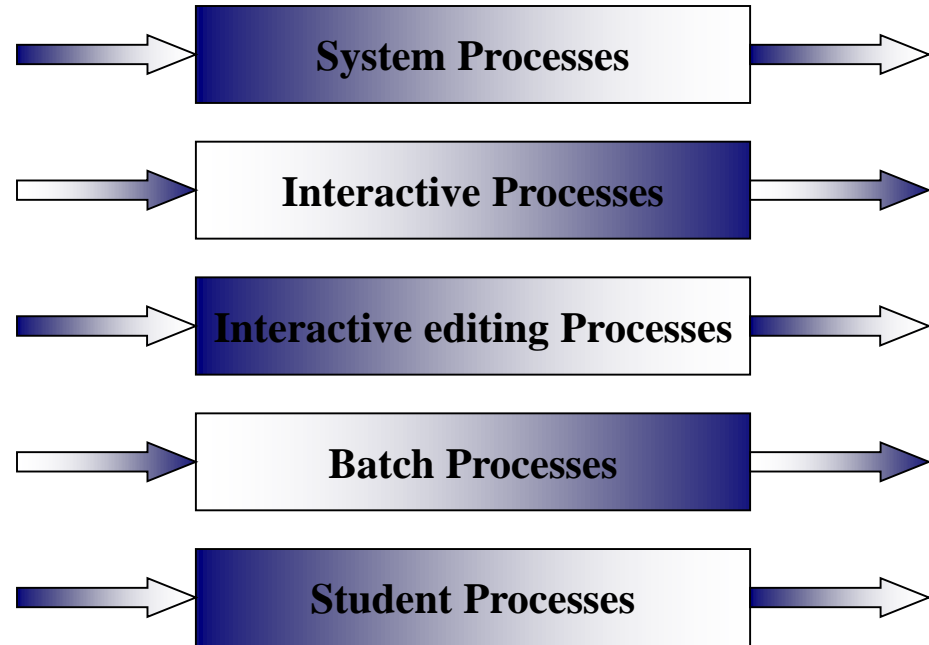
# Multilevel Queue Scheduling

## ■ Example:



# Multilevel Queue Scheduling

- Each queue may have absolute priority over the other queue
- Alternatively, Time slice between the queues
  - Time slots can be equal
  - Or
    - 80% time for Foreground processes
    - 20% time for Background processes





# Multilevel Feedback Queue Scheduling

- In Multilevel Queue a process is permanently assigned to a queue
- The queue to which a process should belong is decided statically
- Multilevel Feedback Queue Scheduling:
  - A Process may move between the Queues
  - Aging can be implemented this way.





# Multilevel Feedback Queue Scheduling

- Multilevel-feedback-queue scheduler defined by:
  - Number of queues
  - Scheduling algorithms for each queue
  - Method used to select when upgrade process
  - Method used to select when demote process
  - Method used to determine which queue a process will enter when that process needs service



# Multilevel Feedback Queue Scheduling

## ■ Example

- If a process used too much CPU time, then move it to a lower-priority queue
- If a process waits too long in a lower priority queue, then move it to a higher priority queue

# Multilevel Feedback Queue Scheduling

- Example: Three queues:

- $Q_0$  – RR time quantum 8 milliseconds
- $Q_1$  – RR time quantum 16 milliseconds
- $Q_2$  – FCFS

- Scheduling

- A new job enters queue  $Q_0$  served by RR.
- Then job receives 8 milliseconds.
- If not finished in 8 milliseconds, moved to  $Q_1$ .
- At  $Q_1$  job served by RR.
- Then receives 16 milliseconds.
- If not complete, preempted and moved to  $Q_2$ .

