

Lecture 5

Threads

Operating Systems



Thread: Introduction

- Each process has
 1. Own Address Space
 2. Single thread of control
- A process model has two concepts:
 1. Resource grouping
 2. Execution
- Sometimes it is useful to separate them



Unit of Resource Ownership

- A process has an
 - Address space
 - Open files
 - Child processes
 - Accounting information
 - Signal handlers
 - Etc
- If these are put together in a form of a process, can be managed more easily

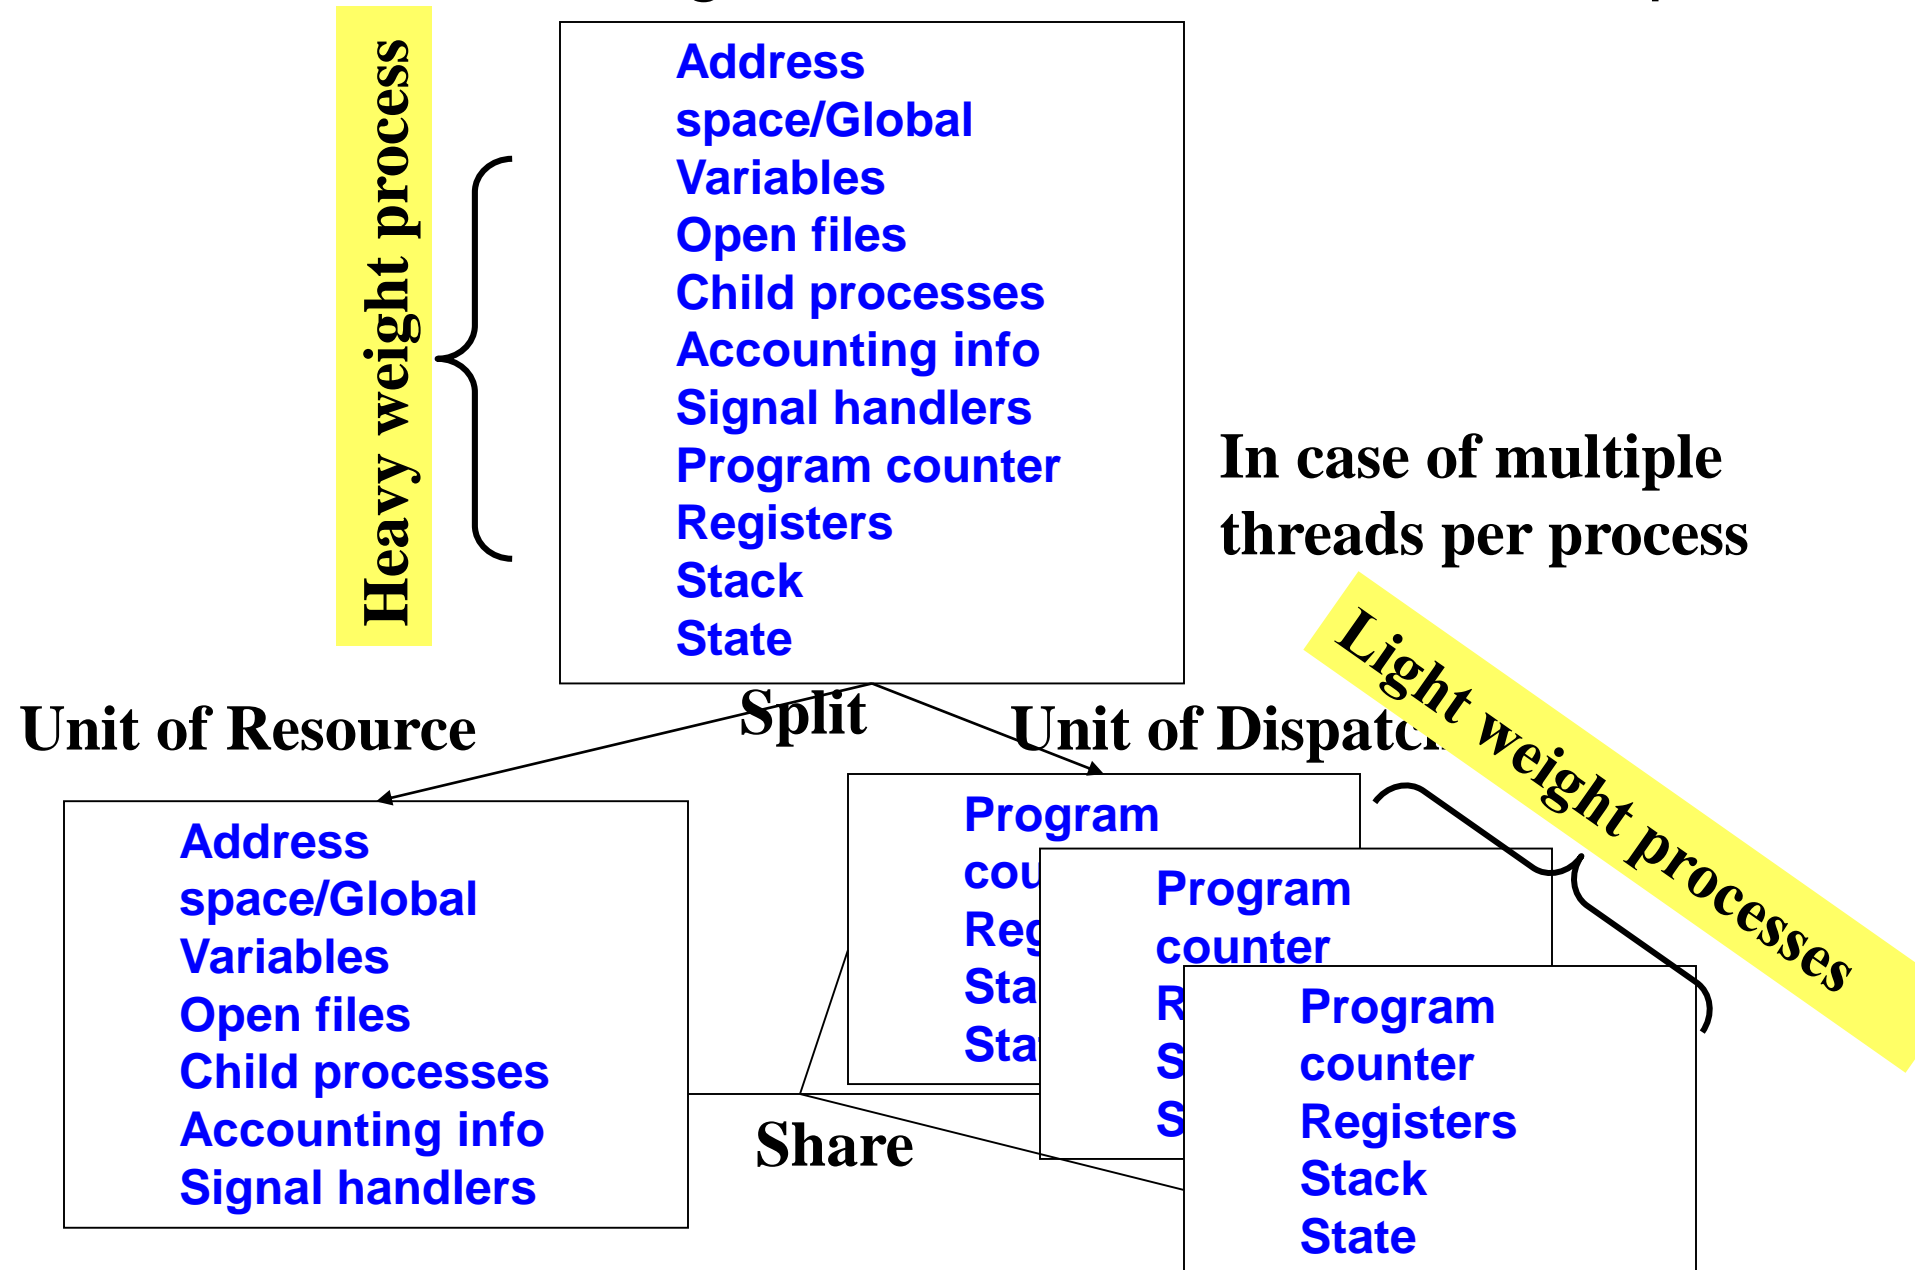
Unit of Dispatching

■ Path of execution

- Program counter: which instruction is running
- Registers:
 - holds current working variables
- Stack:
 - Contains the execution history, with one entry for each procedure called but not yet returned
- State

- Processes are used to group resources together
- Threads are the entities scheduled for execution on the CPU
- Threads are also called ***lightweight*** process


Its better to distinguish between the two concepts



Threads

- Allow multiple execution paths in the same process environment
- Threads share address space, open files etc
- But have own set of Program counter, Stack etc
- The first thread starts execution with `int main(int argc, char *argv[])`
- The threads appear to the Scheduling part of an OS just like any other process


Process Vs. Threads



(a) Three threads, each running in a separate address space

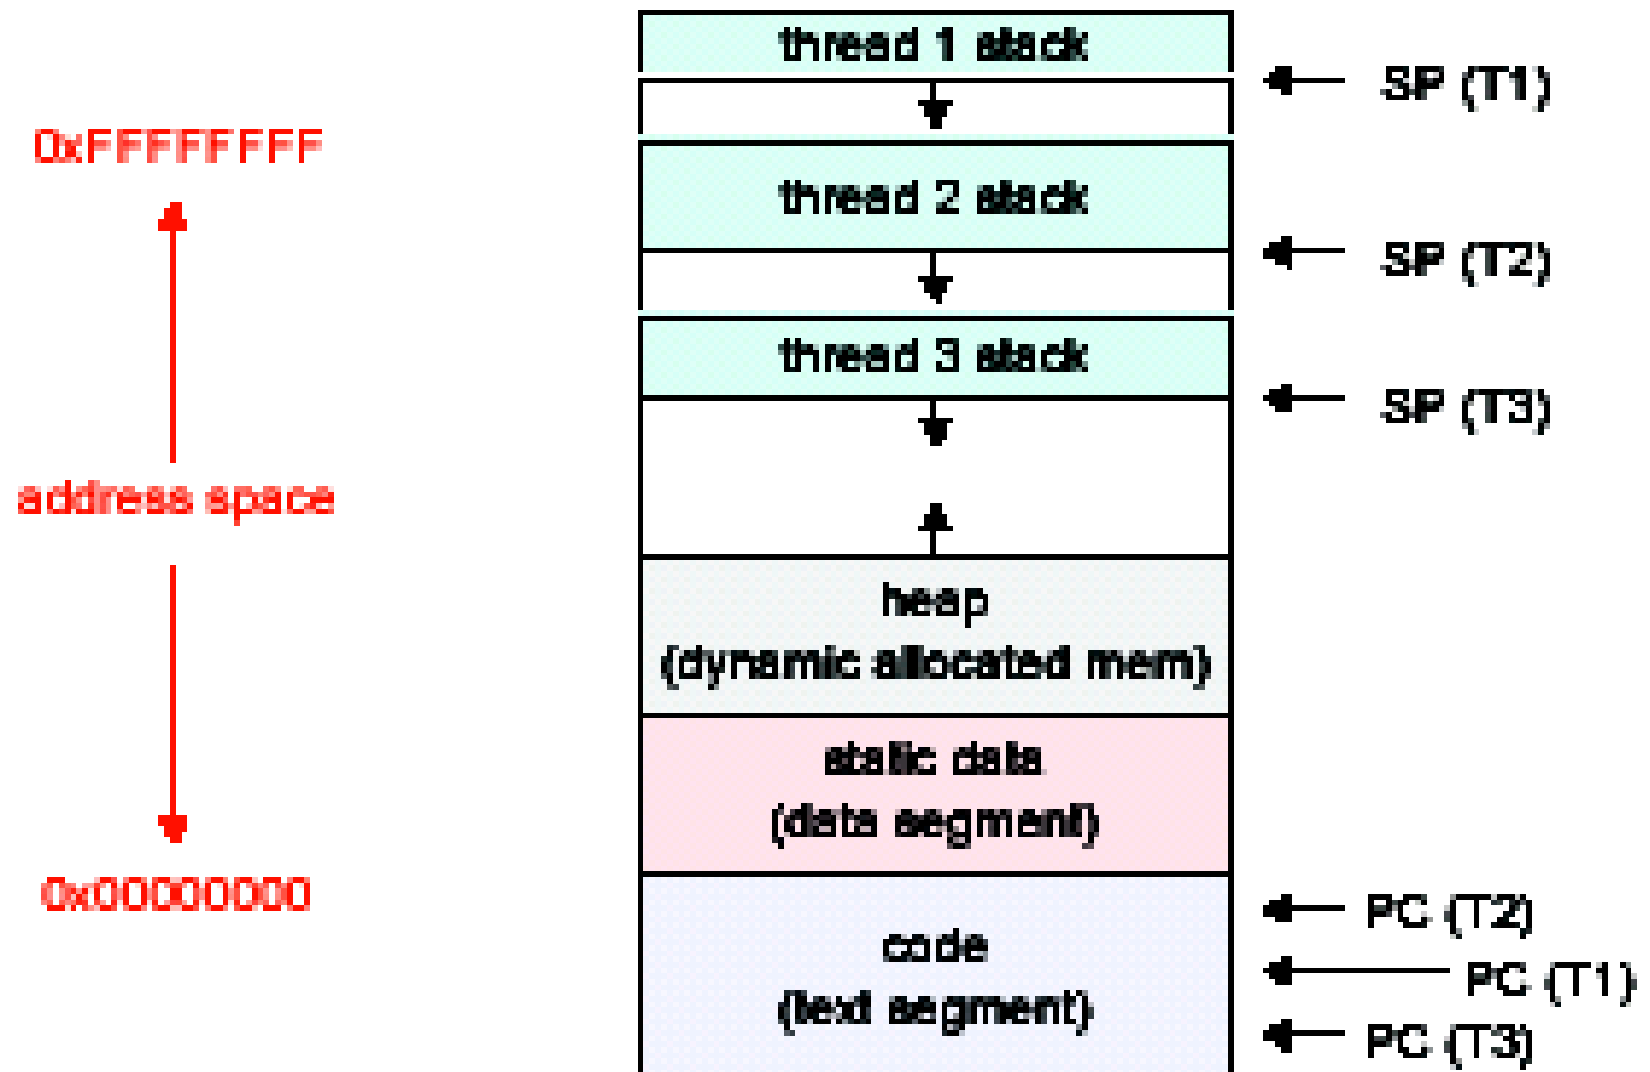
(b) Three threads, sharing the same address space

The Thread Model



Each thread has its own stack

(new) Address space with threads






Thread Usage

- Less time to create a new thread than a process
 - the newly created thread uses the current process address space
 - no resources attached to them
- Less time to terminate a thread than a process.
- Less time to switch between two threads within the same process, because the newly created thread uses the current process address space.
- Less communication overheads
 - threads share everything: address space, in particular. So, data produced by one thread is immediately available to all the other threads
- Performance gain
 - Substantial Computing and Substantial Input/Output
- Useful on systems with multiple processors


Single threaded database server

- Handles multiple clients
 - Either handle the requests sequentially
 - Or multiplex explicitly by creating multiple processes
- Problem with 1
 - Unfair for quick requests, occurring behind lengthy requests
- Problem with 2
 - Complex and error prone
 - Heavy IPC required

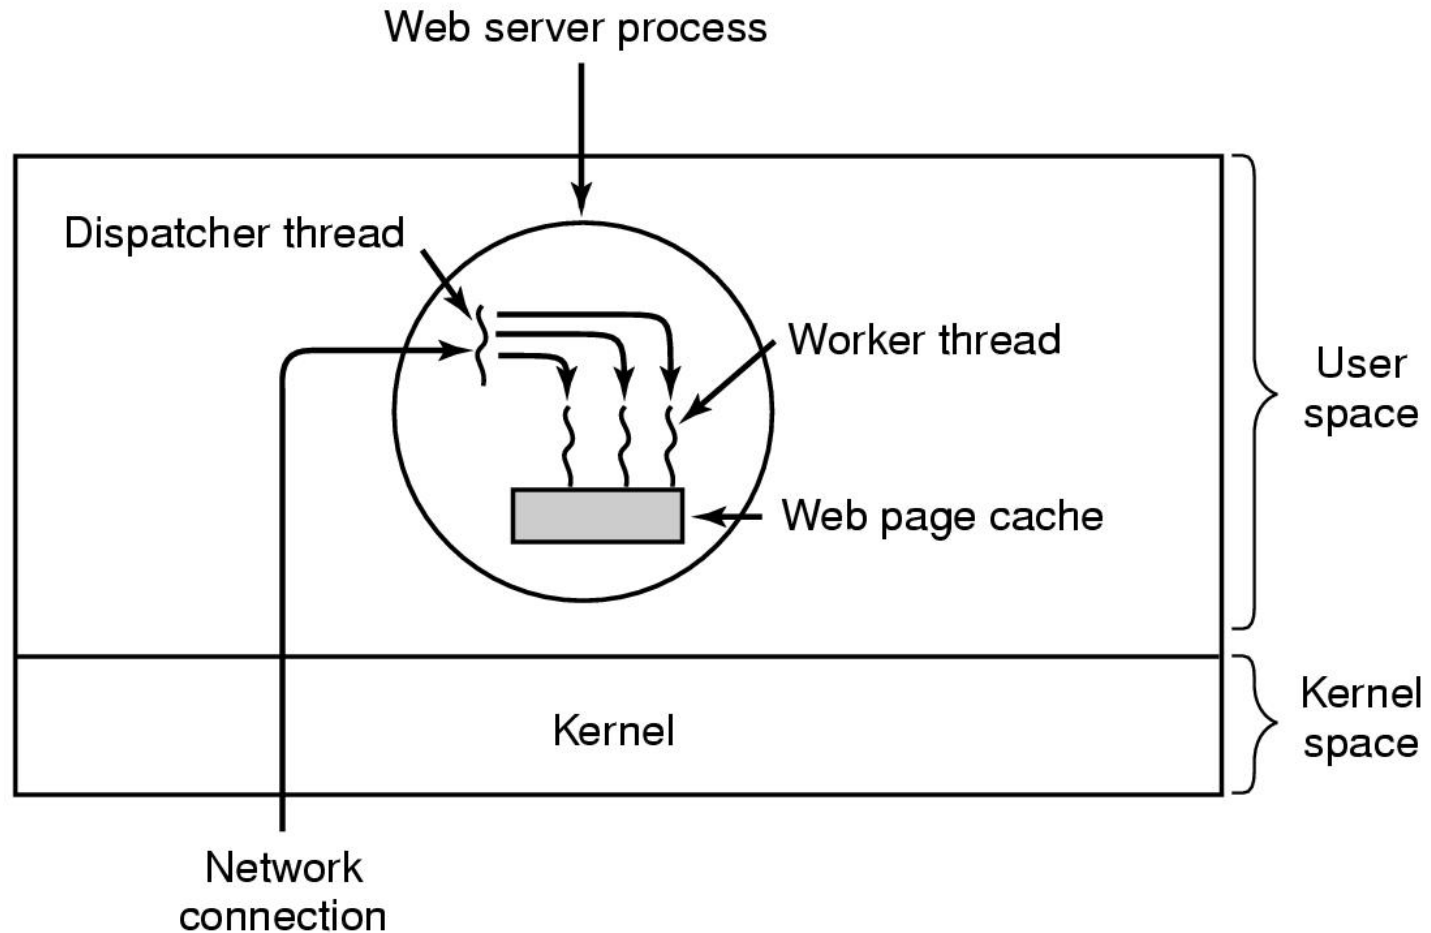


1. Multithreaded database server

- Assign a separate thread to each request
- As fair as in the multiplexed approach.
- The code is as simple as in the sequential approach, since the address space is shared => all variables are available
- Some synchronization of access to the database is required, this is not terribly complicated.



e.g. A multithreaded web server




2. Background Processing


- Consider writing a GUI-based application that uses:
 - Mouse
 - Keyboard input
 - Handles various clock-based events
- In a single threaded application, if the application is busy with one activity, it cannot respond (quickly enough) to other events, such as mouse or keyboard input.
- Handling such concurrency is difficult and complex
- But simple in a multithreaded process



e.g. A word processor with 3 threads



3. Parallel Algorithms e.g. Merge Sort



- Sort some data items on a shared-memory parallel computer.
- Our task is merely to implement a multithreaded sorting algorithm.
 - Divide the data into four pieces
 - Have each processor sort a different piece.
 - Two processors each merge two pieces
 - One processor merges the final two combined pieces.



Application benefits of threads

- Consider an application that consists of several independent parts that do not need to run in sequence
- Each part can be implemented as a thread
- Whenever one thread is blocked waiting for an I/O, execution could possibly switch to another thread of the same application (instead of switching to another process)



Benefits of Threads

- Since threads within the same process share memory and files, they can communicate with each other without invoking the kernel
- Therefore necessary to synchronize the activities of various threads so that they do not obtain inconsistent views of the data

Example of inconsistent view

- 3 variables: A, B, C which are shared by thread T1 and thread T2
- T1 computes $C = A + B$
- T2 transfers amount X from A to B
 - T2 must do: $A = A - X$ and $B = B + X$ (so that $A + B$ is unchanged)
- But if T1 computes $A + B$ after T2 has done $A = A - X$ but before $B = B + X$
- then T1 will not obtain the correct result for $C = A + B$