

Lecture 6


Introduction to

Memory Management

Operating Systems

Memory Management

- A programmer will like to have memory
 - Infinitely large
 - Infinitely fast
 - Non volatile
- However, we have Memory Hierarchy:



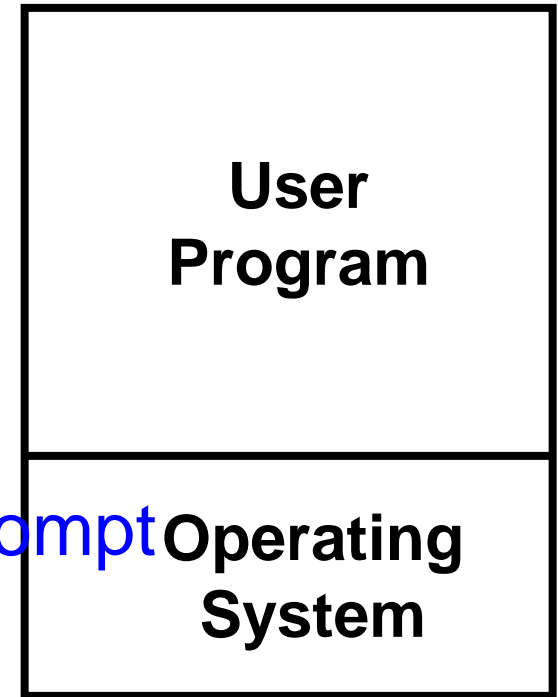


Memory Manager

- The part of the OS that manages the Memory Hierarchy
 - Which part of memory is in use and which is not in use
 - Allocate memory
 - Deallocate memory
 - Swapping between main memory and disks

Uniprogramming

- Run just one program at a time
- Memory is shared between
 - A User program
 - Operating system
- The user types a command on the prompt
- The OS
 - Copies the program from the disk to memory
 - Executes the program
 - After execution, prompt is available again
- The new program is copied into RAM, overwriting the previous one



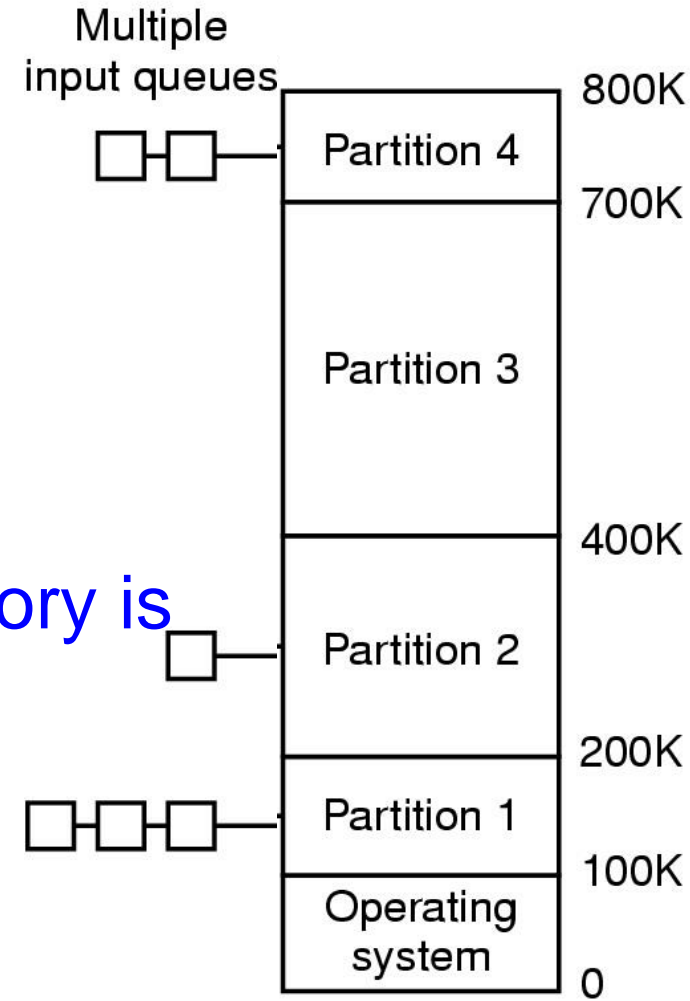
Multiprogramming with Fixed Partitions

■ Multiprogramming:

- When one process blocks due to I/O
- Another one can use the CPU

■ Divide memory into n (usually unequal fixed partitions)


- If the job size < Partition size, memory is wasted
- When a job arrives
 - put it into the queue of



The *smallest partition large enough* to hold the job

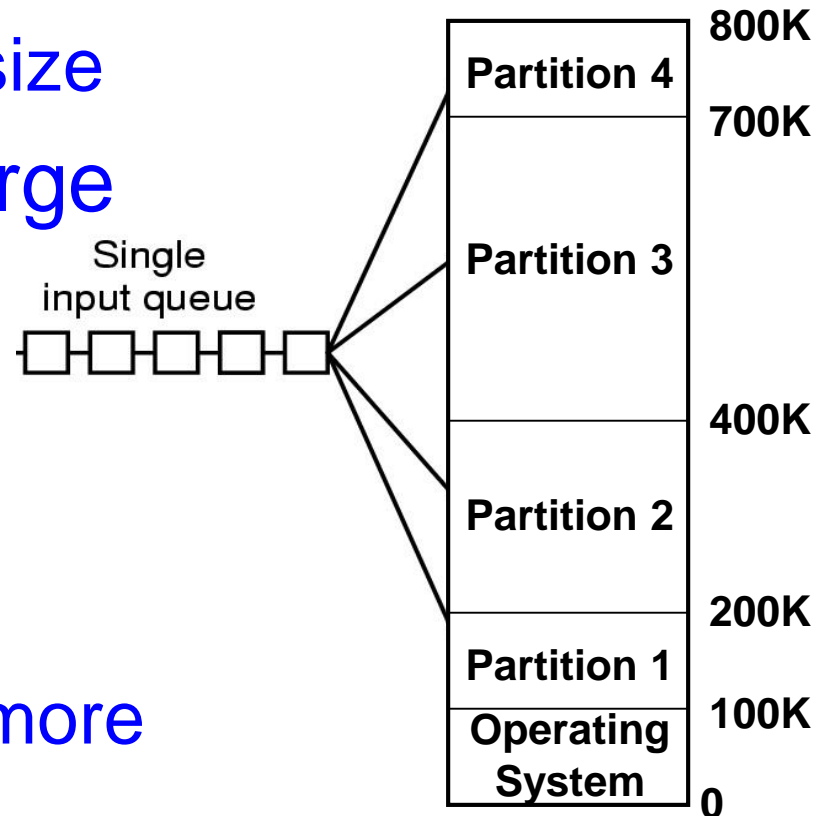
Multiprogramming with Fixed Partitions

- Disadvantage of Multiple Queues:
 - When a large partition is empty
 - And queues for small partition is full
 - Small jobs have to wait, even though plenty of memory is free
- Alternative: Maintain a single Queue




Multiprogramming with Fixed Partitions

- Whenever a partition becomes free, a job is selected
 - Closest to the front of the queue
 - Smaller than the partition size
- Undesirable to waste a large partition for smaller job
 - Search the queue, find the largest job that fits in
- Unfair for smaller jobs
 - A job may not be skipped more than k times




Relocation and Protection

- Different jobs will run at different addresses
- Linker:
 - Combines the user procedures and the library procedures
 - Produces a single binary file
- Suppose, the first instruction is
 - call a procedure at absolute address 10 within the binary file
- If the process is loaded in 1st partition
- This call will jump inside the Operating system



Relocation and Protection


- Solution:
- If the program is loaded in 1st partition, then
 - Call $100k + 10$
- If the program is loaded in 2nd partition, then
 - Call $200k + 10$
- So on...
- This is called **Relocation** problem



Relocation and Protection

■ Solution:

- Add an offset to each address in the program
- The offset depends on the partition
- E.g. if the Program is loaded in partition 1, add 100k to every address
- if the Program is loaded in partition 2, add 200k to every address



- A program can still generate an address that jumps in the OS or other user's code
- We need to **Protect** the code

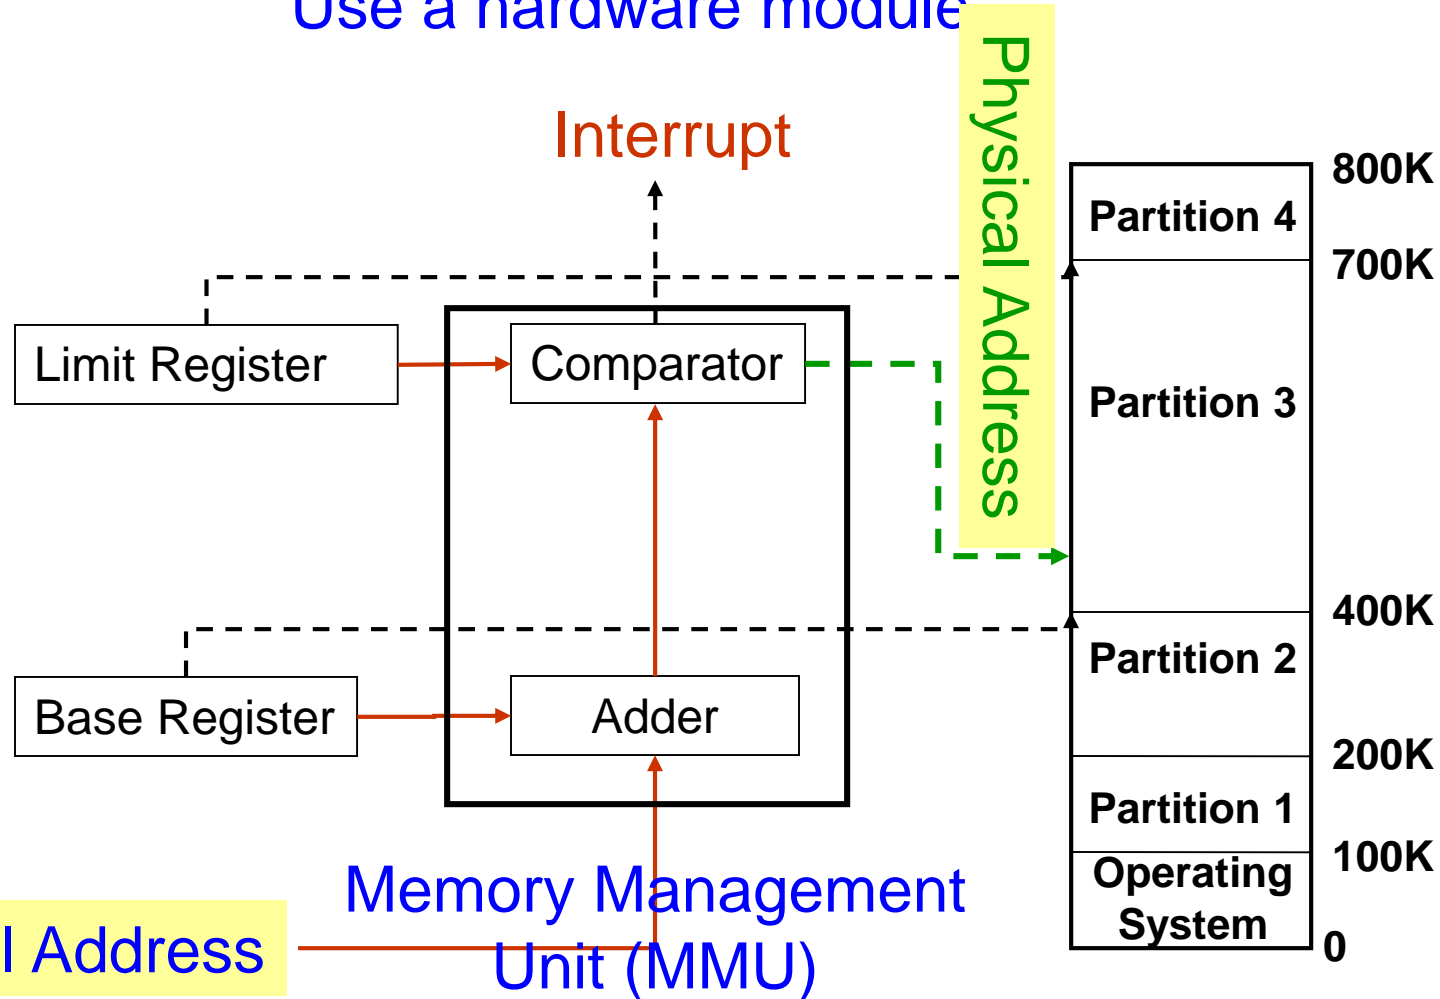
Relocation and Protection

- Solution:
 - Base and Limit Registers
- When a process is scheduled
- Base Register
 - Loaded with the starting address of the Partition
- Limit Register
 - Loaded with the length of the Partition
- Before referring to memory
 - Add the base register contents to generated memory address
- Also check against the Limit register for protection

Relocation and Protection

Addition and comparison has to be performed on every address

Use a hardware module




Internal Fragmentation

Memory that is internal to a partition but not being used

Internal Fragmentation

Internal Fragmentation



Swapping



- If *not enough space* in memory for all the currently active processes
- Excess processes must be kept on disk
- Example of Swapping
 - Round Robin Scheduling
- When the quantum expires
 - Swap out the currently running process
 - Swap in another process to freed memory space



Swapping

- Another Example:
 - Priority based scheduling
- If a higher priority process arrives
 - Swap out a lower priority process
 - Swap in the higher priority process
- When the higher priority process exits
 - Swap in the lower priority process

Swapping



The number, size and location of the partition is decided dynamically.



D is of Higher Priority, A is of Lower Priority!!!

B exits now, Swap in A

Swapping

Its possible to combine all the holes into one big hole

This is called **COMPACTION**



**EXTERNAL
FRAGMENTATION**

Memory External to all partitions is Fragmented

Enough total memory exists to satisfy the request

But is *not Contiguous*

Swapping has created Holes

D exits now

No Space for E !!!

Now E has enough space

Swapping


- Compaction involves CPU overhead
- If addresses are not generated relative to the partition location:
 - Then Compaction is not possible
- How much memory should be allocated for a process?
 - If all the memory is allocated statically in a program
 - Then, OS knows exactly the amount of memory to be allocated: **executable code + variables**
- However, if memory is allocated dynamically (say using **new**)

Swapping

- Problem may occur whenever a process tries to grow
- If a **hole** is adjacent to the growing process then it can be allowed to grow
- If **another process** is adjacent to the growing process, then
 - The growing process should be **moved to a larger hole**
 - If a larger hole is **not** available, then, one or more processes should be **swapped out**
 - If a process **cannot** be swapped out (say, there is not enough space on disk)
 - The growing process has to **wait**
 - Or should be **killed!!!**

Swapping

- If most of the processes grow as they run,
- It is better to allocate a little extra memory for a process



Swapping

- If processes have two growing segments
 - Data (as heap for dynamically allocated variables)
 - Stack

• The memory can be used by either of the two segments

• If it runs out the process either

- Has to be moved to a larger hole
- Or swapped out
- Or Killed

