

Lecture 7

Memory Management

II

Operating Systems




How to keep track of Memory

- We need to keep track of
 - Allocated memory
 - Free memory
- Memory management with bitmaps
- Memory management with linked lists

Memory management with bitmaps

- Memory is divided up into allocation units
 - Few words
 - Or several kilobytes





Memory management with bitmaps


- Size of allocation unit is important
- Smaller allocation unit
 - Larger bitmap required
- Larger allocation unit
 - Smaller bitmap required
 - More memory will be wasted
 - If the process size is not an exact multiple of the allocation unit

Memory management with bitmaps

- To bring a k unit process in memory
- Search for a k run consecutive 0 bits in the map
- Search can be slow
- Since, k run may cross word boundaries

Find run of length = 3


Find run of length = 4



Memory management with Linked Lists

■ Linked list of memory segments


- Free segments → Holes
- Allocated segments → Processes




Memory management with Linked Lists

- Segment list is sorted by addresses
- Sorting helps in updating the list, when a process is swapped out or exits
- A process usually has two neighbors


Updating the list requires




Replace P by H



Two entries are merged



Two entries are merged



Three entries are merged

Memory Allocation with Linked Lists

- Several algorithms for allocating memory with linked list
- FIRST FIT:
 - Scan the segment list, until
 - A hole large enough is found
 - Split the hole into two pieces (if not exact match)
 - One for the new process
 - One for the unused memory
- The algorithm is fast since it searches as little as possible



Memory Allocation with Linked Lists

■ NEXT FIT

- ☐ Works the same as First Fit, except,
 - ☐ Does not always start searching from the beginning
 - ☐ Rather starts searching the list, from where it left last time
- Simulations show, that gives no better performance than First Fit



Memory Allocation with Linked Lists

■ BEST FIT

- Search the entire list
- Take the smallest hole that is adequate

■ Best Fit tries to find the hole closest to the size

■ Slower than First Fit

- Every time has to search the entire list

■ But still results in more memory wastage

- Tends to fill up memory, with tiny useless holes
- First Fit generates larger holes on the average

Memory Allocation with Linked Lists

■ WORST FIT


- Take the largest hole available
- So that, the hole broken off will be large enough to be useful
- What if the larger holes left by worst fit are not useful
- → more memory wasted than Best Fit

Memory Allocation with Linked Lists

- Search time of all the four algorithms can be improved by
- Keeping separate lists for Process and Holes
- While allocating memory only have to search the lists of holes
- DRAWBACK:
 - Problem while deallocating memory
 - A node from the process list has to be inserted in the hole list

Memory Allocation with Linked Lists

- The Hole list can be kept sorted by size
- As soon as a hole that fits is found, no more searching is required
- Another improvement can be,
 - Maintain different lists for different sizes of holes
 - QUICK FIT



Memory Allocation with Linked Lists

- QUICK FIT

- Finding a hole of required size is extremely fast

- DRAWBACK:

- ☐ When a process terminates/swapped out
- ☐ Finding the neighbor to see whether the merge is possible
- ☐ If merge is not done, then sooner memory will be fragmented into large holes